
1 Overview

- 1.1 Location** `$<AMDAPPSDKSamplesInstallPath>\samples\bolt\`
- 1.2 How to Run** See the *Getting Started* guide for how to build samples. You must first compile the sample. Use the command line to change to the directory where the executable is located. The pre-compiled sample executable is at `$<AMDAPPSDKSamplesInstallPath>\samples\bolt\bin\x86\` for 32-bit builds, and `$<AMDAPPSDKSamplesInstallPath>\samples\bolt\bin\x86_64\` for 64-bit builds.
- At the command prompt, type `BoltIntro` to run the program.

2 Introduction

Bolt is a C++ template library optimized for GPUs. Bolt provides high-performance library implementations for common algorithms such as scan, reduce, transform, and sort. The Bolt interface resembles the C++ Standard Template Library (STL); developers familiar with STL will recognize many of the Bolt APIs and customization techniques. In some cases, developers can benefit from GPU acceleration simply by changing the namespace for STL algorithm calls from `std` to `bolt::cl`.

Bolt can directly interface with host memory structures such as `std::vector` or host arrays (`float*`). On current GPU systems the host memory is automatically mapped or copied to the GPU. Bolt also provides a `bolt::cl::device_vector` container that can be used to allocate and manage device-local memory for higher performance on discrete GPU systems. Bolt APIs can accept both STL container or `device_vector` iterators.

C++ templates can be used to customize the algorithms with new types (for example, the Bolt sort can operate on ints, float, or any custom type defined by the user). Additionally, Bolt lets users customize the template routines using function objects (functors) written in OpenCL; for example, to provide a custom comparison operation for sort, or a custom reduction operation.

3 Implementation Details

The example is very similar to that of the familiar C++ Standard Template Library; the only difference is the include file (`bolt/cl/sort.h`) and the `bolt::cl` namespace before the sort call. If the include is removed and the `bolt::cl` namespace changed to `std`, the code uses the STL version of sort. Bolt developers do not need to learn a new device-specific programming model to leverage the power and performance advantages of GPU computing.

The example demonstrates three important Bolt features.

- The host-allocated vector `randInput` is directly passed to the Bolt sort routine, without a need to explicitly allocate and manage GPU device memory.
- The Bolt `bolt::cl::sort` call submits to the platform, rather than a specific device. The Bolt runtime selects the best device to run the sort, potentially running it on the CPU in the event a GPU is not available or the sort size is too small to benefit from GPU acceleration.
- The Bolt `bolt::cl::sort` call also can work with arrays, such as `arrInp`, created on the host.

User-defined datatypes can also be sorted using `bolt::cl::sort` by declaring the type under `BOLT_FUNCTOR()`, overriding the necessary operator functions and using `BOLT_CREATE_TYPENAME()` to register the type with Bolt (for example: `BOLT_CREATE_TYPENAME(bolt::cl::greater< MyType<int> >)`). The implementation of this is left as an exercise to the user.

4 Requirements

Bolt requires an OpenCL™ implementation that supports the static C++ kernel features; specifically, C++ template support for OpenCL™ kernels. Currently, the AMD runtime provides this support; other vendors may adopt this feature in the future.

5 References

1. <http://developer.amd.com/tools-and-sdks/opencl-zone/bolt-c-template-library/>

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:

URL: developer.amd.com/appsdk
Developing: developer.amd.com/



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks

© 2015 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.