

2D Fluid Simulation Using the Lattice-Boltzmann Method

1 Overview

1.1 Location `$<AMDAPPSDKSamplesInstallPath>\samples\opencl\cl\1.x`

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The default executables are placed in `$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86` for 32-bit builds and `$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86_64\` for 64-bit builds.

Type the following command(s).

1. `FluidSimulation2D`
This runs the simulation in a 256x256 grid.
2. `FluidSimulation2D -h`
This prints the help message.

1.3 Command Line Options Table 1 lists, and briefly describes, the command line options.

Table 1 Command Line Options

Short Form	Long Form	Description
-h	--help	Shows all command options and their respective meaning.
	--device	Devices on which the program is to be run. Acceptable values are <code>cpu</code> or <code>gpu</code> .
-q	--quiet	Quiet mode. Suppresses all text output.
-e	--verify	Verify results against reference implementation.
-t	--timing	Print timing.
	--dump	Dump binary image for all devices.
	--load	Load binary image, and execute on device.
	--flags	Specify compiler flags to build kernel.
-p	--platformId	Select platformId to be used (0 to N-1, where N is the number of available platforms).
-d	--deviceId	Select device to be used (0 to N-1, where N is the number of available devices).
-v	--version	AMD APP SDK version string.
-i	--iterations	Number of iterations for kernel execution.

2 Introduction

This sample is based on the Lattice-Boltzmann method (LBM) for simulating a fluid on a 2D grid. The format is D2Q9, which represents the second dimension and nine frequency distributions of velocity. Each cell has nine velocity directions, with indexes shown in Figure 1 (not necessarily in the same order as the sample).




Figure 1 Velocity Directions

The velocity and density values of a cell can be calculated by using the equations

$$\rho = \sum f_i \quad u = \frac{1}{\rho} \sum f_i e_i$$

where

f_i is the frequency of each velocity (the number of particles in a cell going in that particular direction).

i represents values from 0-8, inclusive, denoting a particular velocity, as shown in Figure 1.

e_i denotes the direction vector of the velocity.

The Lattice-Boltzmann equation can be written with respect to space and time as

$$f_i(x, t + \Delta t) - f_i(x, t) = \Omega_i$$

where

Ω_i is a fluid collision operator that models the collision between fluid molecules in a cell.

We use a simple collision operator, called the Bhatnagar-Gross-Krook approximation, which uses a single relaxation time approximation to reduce the operator to operations suitable for computers. It is based on the idea that the main effect of the collision operator is to bring the molecule distribution closer to the equilibrium distribution, which is defined by

$$f_i^{eq} = \omega_i \rho \left(1 - \frac{3}{2} u^2 + 3(e_i \cdot u) + \frac{9}{2} (e_i \cdot u)^2 \right)$$

where

ω_i is a constant that depends on the lattice geometry.

The collision operator is defined as

$$\Omega_i = -\frac{\Delta t}{\tau} \left(f_i(x, t) - f_i^{eq}(p, u) \right)$$

where

τ is a constant that represents the viscosity of the fluid.

Combining the equations

$$f_i(x + e_i \Delta t, t + \Delta t) = f_i(x, t) - \frac{\Delta t}{\tau} \left(f_i(x, t) - f_i^{eq}(p, u) \right)$$

The basic steps of the LBM are:

1. Iteration start – Initialize the values of the velocity distributions for each cell.
2. Collision step – Compute the density and scalar velocity of a cell, and calculate the equivalent distribution.
3. Streaming step – Distribute the newly computed frequency distribution (according to final equation) to neighbors.

The diagram in Figure 2 explains the 3 steps.




Figure 2 Basic LBM Steps

3 Implementation Details

The simulation is initialized and updated by holding the left-mouse button down while dragging the mouse across the entire grid. This assigns velocity values to the cells, and the simulation continues from there.

There are two sets of buffers: one for the input frequency distribution, the other for the output frequency distribution. These store nine values for each cell in the grid. There is one velocity buffer which stores the resultant velocity for each cell. There is also one `type` buffer, which differentiates between empty space and boundary cells. The boundary cells are created dragging the right-mouse button on the grid.

In each step, the kernel updates the output buffer with the new distribution, which becomes the next input. The velocity buffer is translated to RGB values using a color conversion table, and this buffer is displayed using a GL renderer.

The sample uses double precision, so it runs only on following devices: Radeon 4800 series, Radeon 5800 series, and CPU. The color conversion is done on the host; thus, with every step buffers are copied to the host. As a result, the GPU might process fewer frames per second than the CPU.

4 References

1. The OpenCL code is based on the openMP code available at <http://software.intel.com/en-us/courseware/course/view.php?id=87>.
2. The theory presented is taken from <http://www.monitzer.com/FluidGPU/>.

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:
URL: developer.amd.com/appsdk
Developing: developer.amd.com/
Support: developer.amd.com/appsdksupport



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks

© 2015 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.