

1 Overview

1.1 Location `$<AMDAPPSDKSamplesInstallPath>\samples\opencl\cl\1.x`

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The pre-compiled sample executable is at

`$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86\` for 32-bit builds, and
`$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86_64\` for 64-bit builds.

Type the following command(s).

1. `SimpleConvolution`
Performs convolution of a 64x64 image with a blur mask of 3x3.
2. `SimpleConvolution -h`
This prints the help file.

1.3 Command Line Options Table 1 lists, and briefly describes, the command line options.

Table 1 Command Line Options

Short Form	Long Form	Description
-h	--help	Shows all command options and their respective meaning.
	--device	Devices on which the program is to be run. Acceptable values are <code>cpu</code> or <code>gpu</code> .
-q	--quiet	Quiet mode. Suppresses all text output.
-e	--verify	Verify results against reference implementation.
-t	--timing	Print timing.
	--dump	Dump binary image for all devices.
	--load	Load binary image and execute on device.
	--flags	Specify compiler flags to build the kernel.
-p	--platformId	Select platformId to be used (0 to N-1, where N is the number of available platforms).
-d	--deviceId	Select deviceId to be used (0 to N-1, where N is the number of available devices).
-v	--version	AMD APP SDK version string.
-x	--width	Width of the input matrix.
-y	--height	Height of the input matrix.

Short Form	Long Form	Description
-m	--masksize	Dimension of convolution filter.
-i	--iterations	Number of iterations for kernel execution.
-l	--localSize	Work-group Size.

2 Introduction

Convolution or image filtering is widely used in Image Processing applications such as blur, smooth effects, or edge detection. The sample shows the naive implementation of convolution. The sample supports only 3X3 and 5X5 Sobel Filter.

The convolution or filtering operation can be implemented in two ways. One is a separable convolution, meaning that the filtering or convolution is done in two separate passes. The other is a non-separable convolution, where convolution is done in one pass. The advantage of separable filter over the non-separable filter is that it reduces the number of arithmetic operation. For example, naïve convolution with 2-D square filter of size ($m \times m$) requires m^2 number of mad (multiplication and addition) operations for each pixel, whereas separable version of same filter reduces the number to $2m$.

OpenCL implementation of fixed sized separable convolution with 2 passes could be done in several ways. In this sample, we demonstrate the approach of using 2 separate kernels for each of the passes. A more advanced version of separable convolution using a single kernel for both the passes is demonstrated in the AdvancedConvolution SDK sample.

3 Implementation Details

A convolution filter is just a scalar product of the filter weights with the input pixels within a window surrounding each of the output pixels.

This sample includes two kernels:

- **Non-separable Filter – [One pass: a naïve implementation]:**
Non-separable filter is a one pass convolution operation. This kernel implements 2D convolution on device global memory. In non-separable convolution, each pixel of the output image is calculated by fetching raw pixels, performing a scalar product of each pixel with pre-defined filter coefficients.
- **Separable Filter - [Two passes]:**
This kernel implements a two pass convolution of a 2D separable filter. It uses 2 separate kernels; one is horizontal filter (or row-wise filter) and second is vertical filter (or column-wise filter). The first pass applies a one-dimensional horizontal filter to the entire image row of pixels, does additional work of transposing the output and store result column-wise rather than row-wise. In second pass, although it works along the row as well, actually it does vertical filtering, since the frame has already been transposed in first pass and finally it transposes the output so that resulting image is restored in its original orientation.

4 References

1. <http://en.wikipedia.org/wiki/Convolution>
2. http://en.wikipedia.org/wiki/Separable_filter
3. <http://blogs.mathworks.com/steve/2006/10/04/separable-convolution/>

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:
URL: developer.amd.com/appsdk
Developing: developer.amd.com/
Support: developer.amd.com/appsdksupport



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks

© 2015 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.