

Region Growing Segmentation

1 Overview

1.1 Location `$<AMDAPPSDKSamplesInstallPath>\samples\opencl\cl\2.0`

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The pre-compiled sample executable is at

`$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86\` for 32-bit builds, and
`$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86_64\` for 64-bit builds. Ensure that the OpenCL 2.0 environment is installed.

Type the following command(s).

1. `RegionGrowingSegmentation`
This command runs the program with the default options.
2. `RegionGrowingSegmentation -h`
This command prints the help file.

1.3 Command Line Options Table 1 lists, and briefly describes, the command line options.

Table 1 Command Line Options

Short Form	Long Form	Description
-h	--help	Shows all command options and their respective meanings.
	--device [cpu gpu]	Devices on which the OpenCL kernel is to be run. Acceptable values are <code>cpu</code> or <code>gpu</code> .
-q	--quiet	Quiet mode. Suppresses all text output.
-e	--verify	Verify results against reference implementation.
-t	--timing	Print timing-related statistics.
-v	--version	AMD APP SDK version string.
	--dump [filename]	Dump the binary image for all devices.
	--load [filename]	Load the binary image and execute on the device.
	--flags [filename]	Specify the filename containing the compiler flags for building the kernel.
-i	--iterations	Number of iterations.
-p	--platformId	Select the platformId to be used[0 to N-1 where N is number platform s available].
-d	--deviceId	Select deviceId to be used[0 to N-1 where N is number devices available].

Short Form	Long Form	Description
-o	--hostq	Host side enqueue of kernels.

2 Introduction

This sample shows the use of the device-side enqueue feature in OpenCL 2.0, in scheduling dynamically generated work by a device kernel on the device itself. This dynamic generation of work and re-scheduling it on the device is shown by an example of seeded region growing image segmentation. Given an image, the user provides seed pixels, each inside one region to be segmented out. The seed pixels are put in a queue. Starting with the seed pixels, the algorithm classifies already queued pixels into one of the regions, which is chosen based on minimum difference between pixel luma and average luma of the region. It then queues the unclassified neighbors of the queued pixels to be classified in the next iteration. With each iteration average luma of all the regions are also updated.

The algorithm is adapted from seeded region growing algorithm of Adams and Bischof[1].

This sample must be run in the OpenCL 2.0 environment.

3 Implementation

The host reads an input image “input.bmp”, and a seed file called “seeds.txt”. The top row of the seed file contains the number of regions into which the image is to be segmented, and rest of the rows contains seed pixels co-ordinates, each for one region. The images should be in 24 bit RGB bitmap format.

The sample uses device side enqueue feature to do image segmentation. This feature is used to repeatedly enqueue dynamically generated work on device as follows.

1. In the beginning, all seeded pixels are classified to their respective regions and put in a queue by an initialization kernel.
2. A top level kernel named `grow_region` then evaluates amount of work to be done based on the queue size and enqueues another kernel `queue_neigh` having work items equal to size of the queue.
3. `queue_neigh` finds unclassified pixels in the neighborhood of each queued pixel and forms another queue. It then enqueues a kernel `classify` to classify pixels in the newly formed queue.
4. `classify` classifies pixels in the new queue to regions, and again enqueues `grow_region`. `grow_region` dynamically evaluates the work to be done based on the size of the new queue and enqueues `queue_neigh`.

The iteration continues till no pixels remains to be queued.

The output of device side enqueue is written in “ocl_output.bmp”. a host side verification code writes output in “ver_output.bmp”.

4 References

1. Adams R. and Bischof L., "Seeded Region Growing", IEEE Trans on PAMI, V.16, NO.6, pp 641-647.
2. The OpenCL Specification (ver 2.0, rev 22) document.
3. The OpenCL C Programming Language (ver 2.0, rev 22) document.

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:
URL: developer.amd.com/appsdk
Developing: developer.amd.com/
Support: developer.amd.com/appsdksupport



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks

© 2015 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.