

SVM Atomics Binary Tree Insert

1 Overview

1.1 Location `$<AMDAPPSDKSamplesInstallPath>\samples\opencl\cl\2.0`

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The pre-compiled sample executable is at

`$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86\` for 32-bit builds, and
`$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86_64\` for 64-bit builds.

Ensure that the OpenCL 2.0 environment is installed and that the SVM Fine Grain Atomics feature is supported. Currently on AMD platforms, SVM Fine Grain Atomics feature is supported on only Linux 64-bit.

Type the following command(s).

1. `SVMAtomicsBinaryTreeInsert`
This command runs the program with the default options.
2. `SVMAtomicsBinaryTreeInsert -h`
This command prints the help file.

1.3 Command Line Options Table 1 lists, and briefly describes, the command line options.

Table 1 Command Line Options

Short Form	Long Form	Description
-h	--help	Shows all command options and their respective meanings.
	--device [cpu gpu]	Devices on which the OpenCL kernel is to be run. Acceptable values are <code>cpu</code> or <code>gpu</code> .
-q	--quiet	Quiet mode. Suppresses all text output.
-e	--verify	Verify results against reference implementation.
-t	--timing	Print timing-related statistics.
-v	--version	AMD APP SDK version string.
	--dump [filename]	Dump the binary image for all devices.
	--load [filename]	Load the binary image and execute on the device.
	--flags [filename]	Specify the filename containing the compiler flags for building the kernel.
-i	--iterations	Number of iterations.
-p	--platformId	Select the platformId to be used[0 to N-1 where N is number platform s available].

Short Form	Long Form	Description
-d	--deviceId	Select deviceId to be used[0 to N-1 where N is number of devices available].
-s	--seed	Seed to the random number generator.
-n	--numins	Number of nodes to insert (< 125000000)
-in	--initnodes	Number of initial nodes (> 1)
-hp	--host	Percentage of nodes to be inserted on the host.
-po	--prorder	Print order of the tree nodes.
-g	--gui	Display tree insert animation
-np	--numpasses	Number of passes for node insertion (used only in gui mode, else always 1)

2 Introduction

Building a binary tree with a large number of nodes can be a very time consuming operation when executed on the CPU. Using OpenCL on GPUs, one can accelerate this task significantly. However implementing node insertion into a binary tree poses a few challenges with OpenCL 1.x.

1. The typical data structure of a binary tree includes pointers that point reference to left and right child nodes. Passing data structures with pointers is not allowed in OpenCL 1.x platforms. Hence, the only way to work around this in OpenCL 1.x is to transform the pointer based data structure to indices based structure and pass it to OpenCL kernels. Once a kernel operation is complete, the data must be transformed back to pointer based structure to be used on the host. These data transform operations can be very costly, and affect the overall performance. The code complexity also increases due to the additional data transformations.
2. Furthering accelerate the node insertion operation by having the host and OpenCL device simultaneously inserting nodes into the tree is not possible with OpenCL 1.x. This is because, for concurrent node insertion, synchronization between the host and the OpenCL device is required and there is no way to do this in OpenCL 1.x.

OpenCL 2.0 helps overcome both the above limitations. It supports passing pointer-based data structures to kernels and also enables synchronization between the host and the OpenCL device through C++11 Atomics. It helps in accelerating node insertion and at the same time improves programmability.

3 Implementation

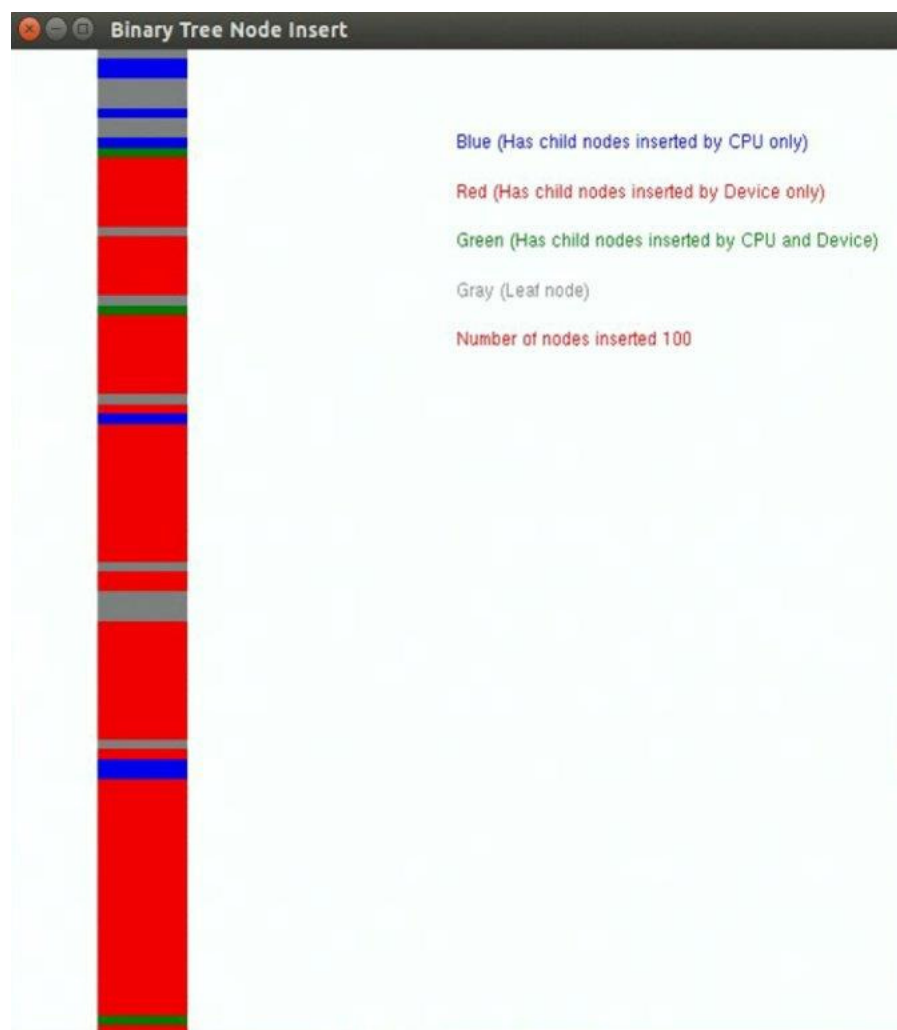
The sample workflow is as follows:

1. Host creates binary tree with a given number of initial nodes.
2. The tree structure is created as a SVM Fine Grain buffer with Atomics flag. This is required for the atomic operations to work.

```
int flags = CL_MEM_READ_WRITE | CL_MEM_SVM_FINE_GRAIN_BUFFER |
CL_MEM_SVM_ATOMICS;
node* svmTreeBuf = (node *) clSVMAlloc(context, flags,
total_nodes*sizeof(node), 0);
```
3. The sample reads the total number of nodes , ' n', to be further inserted into the tree.

4. The 'n' nodes to be inserted can be done all on OpenCL device or all on host or a combination of both. This is decided based on the value of host compute percentage (-hp) passed as input. For example, if -hp value is 0, entire node insertion happens on OpenCL device while if its 100 its completely on host. If -hp is say, 10, host inserts 10% of the total nodes to be inserted and device the remaining 90%.
5. The host code uses OpenMP to parallelize the host part of insertion operation.
6. C++ Atomic operations are used on both the host and the OpenCL kernel to ensure that at any given instance of time, only one thread (either on the host or on the device) inserts to a given parent node.

The sample also has a GUI. It shows nodes in different colors depending on whether it contains child nodes inserted by host only (Blue), device only (Red), both host and device (Green) or a leaf node (Gray).



4 References

1. Memory Model, section 3.3 and Share Virtual Memory, section 5.6 in [OpenCL 2.0 specification](#)
2. Atomic functions, section 6.13.11, in [OpenCL 2.0 C specification](#)
3. Atomic Weapons: The C++ Memory Model and Modern Hardware. Herb Sutter.
<http://herbsutter.com/2013/02/11/atomic-weapons-the-c-memory-model-and-modern-hardware>

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:
URL: developer.amd.com/appsdk
Developing: developer.amd.com/
Support: developer.amd.com/appsdksupport



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks

© 2015 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.