

Asynchronous Data Transfer

1 Overview

1.1 Location `$<AMDAPPSDKSamplesInstallPath>\samples\opencl\cpp_cl\1.x`

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The pre-compiled sample executable is at

`$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86\` for 32-bit builds, and
`$<AMDAPPSDKSamplesInstallPath>\samples\opencl\bin\x86_64\` for 64-bit builds.

Type the following command(s).

1. `AsyncDataTransfer`
This command runs the program with the default options.
2. `AsyncDataTransfer -h`
This command prints the help file.

1.3 Command Line Options Table 1 lists, and briefly describes, the command line options.

Table 1 Command Line Options

| Short Form | Long Form | Description |
|------------|--------------|---|
| -h | --help | Shows all command options and their respective meanings. |
| | --device | Devices on which the program is to be run. Acceptable values are <code>cpu</code> or <code>gpu</code> . |
| -q | --quiet | Quiet mode. Suppresses all text output. |
| -e | --verify | Verify results against reference implementation. |
| -t | --timing | Print timing-related statistics. |
| | --dump | Dump binary image for all devices. |
| | --load | Load binary image and execute on device. |
| | --flags | Specify compiler flags to build the kernel with. |
| -p | --platformId | Select platformId to be used (0 to N-1, where N is the number of available platforms). |
| -v | --version | AMD APP SDK version string. |
| -d | --deviceId | Select deviceId to be used (0 to N-1, where N is the number of available devices). |
| -i | --iterations | Number of iterations. |
| -x | --size | Size of the input buffer per kernel (in Bytes). |
| -k | --kernels | Number of kernels for execution. |

2 Introduction

The Asynchronous Data Transfer sample demonstrates how to harness asynchronous memory transfer in OpenCL. The sample shows the overlap of kernel execution and data transfer in a device in which multiple same or different kernels are executing. The program workflow consists of three steps:

1. Data transfer (input) from Host to Device
2. Kernel Execution
3. Data transfer (output) from Device to Host

The data transfer step is interleaved with kernel execution step, thus improving overall performance.

The sample includes a sequential (synchronous) implementation in which the kernel execution and data transfer steps are invoked sequentially. The performance of the sequential version is compared with that of the Asynchronous version.

This sample is similar to the TransferOverlap sample. The difference is that the TransferOverlap sample used a single queue for data transfer and kernel execution, while this sample uses separate command queues for data transfer and kernel execution.


3 Implementation

The sample implements two versions of kernel executions:

1. Sequential (synchronous)
2. Asynchronous

3.1 Sequential (synchronous) execution implementation

As the name says, the kernels and data transfers run sequentially (writes-> executions-> reads). To make the operations sequential, one can either use blocking operations (i.e. each call returns after the operation is finished) or attach a wait event list. In this sample, blocking calls are used. The synchronous data transfer scenario is shown below:




The following operations are executed in the order given. T1..T6 represent time instances.

1. T1: Write 1 - Writes data from HOST to DEVICE corresponding to Kernel 1.
2. T2: Work 1 - Kernel-1 executes.
3. T3: Read 1 - Reads data from DEVICE to HOST corresponding to Kernel 1.
4. T4: Write 2 - Writes data from HOST to DEVICE corresponding to Kernel 2.

5. T5: Work 2 - Kernel-2 executes.
6. T6: Read 2 - Reads data from DEVICE to HOST corresponding to Kernel 2.

3.2 Asynchronous execution implementation

In this version, non-blocking operations (write, read, execution) are used for making sure that the operations are asynchronous. Since the operations on each individual kernel are sequential (i.e. first write, then execute, and then read), wait and finish events are attached on each operation. The asynchronous data transfer scenario is shown below:



The following operations are executed in the order given.

1. T1: Work 1 - Writes data from HOST to Device corresponding to kernel 1.
2. T2: Kernel-1 execution starts. This overlaps with operation Write 2, viz. data transfer from HOST to Device for Kernel 2.
3. T3: While Read 1 operation starts for Kernel 1, Kernel-2 executes overlaps.
4. T4: Read 2 operation for Kernel 2.

The snapshot below shows the output of the Application Trace profile done using CodeXL. The overlapping bars show the asynchronous nature of the data transfer and kernel execution.

| | |
|---|------------------------|
| OpenCL | |
| Context 1 (0x0000000000024E300) | |
| Queue 1 - Pitsairm (0x0000000003967640) | |
| Data Transfer | 1024.0 KB READ_BUFFER |
| Kernel Execution | |
| Queue 2 - Pitsairm (0x0000000003967A90) | |
| Data Transfer | 1024.0 KB WRITE_BUFFER |
| Kernel Execution | |
| Queue 0 - Pitsairm (0x00000000039672B0) | |
| Data Transfer | |
| Kernel Execution | work |

While Queue 0 is busy in kernel execution, Queue 1 and Queue 2 are concurrently performing read and write data transfers.

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:
URL: developer.amd.com/appsdk
Developing: developer.amd.com/
Support: developer.amd.com/appsdksupport



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks

© 2015 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.